

To: Business Operations Director

From: Aditya K Nagori

Subject: Demand Forecast


Date: 30/04/2023

Problem Statement: To create a forecast of future demand in energy consumption of a power company to meet expected demand.

EXECUTIVE SUMMARY

Major Findings

1. Energy consumption from the forecast tells us that consumption always goes up after mid-year i.e. requirement is always high in the months of July, August, and September.
2. The data we have is almost normal, it is not right or left-skewed.
3. It is also seen that the 10 AM energy usage is lost the same for the rest of the hours.
4. High autocorrelation at the lags indicates that data is nonstationary and needs to be addressed.
5. Every year data pattern indicates the presence of level, and seasonality in the data there is no constant trend. But there is dual seasonality.
6. The forecasted data indicates that the demand continues the same pattern and numbers every year.



Recommendations for Action

1. The production should be increased when the months of July, August, and September are near.
2. The production should be decreased according to the graph where it is indicating low usage.
3. The performance of the resources is to be tracked to make sure Energy production is well equipped for the changing demand.
4. Develop a contingency plan to support sudden spikes in the consumption of energy.

Analytical Overview

The dataset has 39456 rows of observations and 3 variables with date, time, and consumption that show instances of consumption every hour of every day from 2015-12-31 00:00 till 2020-06-3 23:00. The date columns are converted into date time by merging them.

The dataset is visualized for exploratory data analysis through the distribution of data, demand over time, demand by day, demand by hours of the week, & demand by an hour of the day, etc.

ACF plot is implemented for the dataset. The resulting plot shows a strong correlation between the time series and its lagged values. The plot has the lag and correlation coefficient. The plot has two horizontal lines, representing the 95% and 99% confidence intervals for the correlation coefficients. Any correlation coefficients that fall outside of these confidence intervals are considered statistically significant.

The time series is later decomposed into components depending on the season, level, and trend. The seasonal component is assumed to be periodic with a fixed period. The dataset is later divided into training and validation for 39456 observations with 39288 as the training set & 168 as the testing set.

The regression/dynamic regression models, Holt Winter's, ARIMA, SARIMA, dual seasonality models, neural network, smoothing models, & naive models are applied to the dataset to get the best results.

Documentation

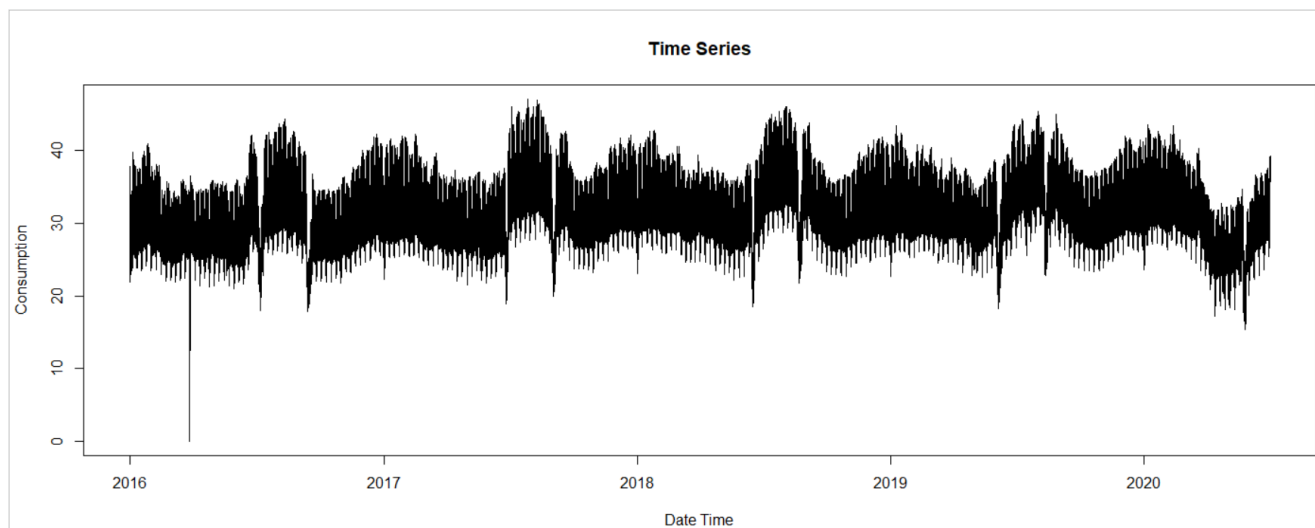
Exploratory Data Analysis:

Getting one column for date time & summarising it.

```
1 setwd("F://Time Series 4th Sem//Project")
2 df= read.csv("Energy.csv")
3 colnames(df) <- c("Date", "Hour", "Consumption (MWH)")
4 df$`Consumption (MWH)` <- gsub(",", "", df$`Consumption (MWH)`)
5 DateTime <- as.POSIXct(paste(df$Date, df$Hour), format="%d.%m.%Y %H:%M")
6 # add the datetime column to your original data frame
7 df <- cbind(df, DateTime)
8 df <- subset(df, select = -c(Date, Hour))
9 # plot the time series using the datetime column
10 plot(df$DateTime, df$`Consumption (MWH)`, type = 'l',
11      xlab = "Date Time", ylab = ("Consumption"), main = "Time Series")
12
13 # Load necessary packages
14 library(ggplot2)
15 library(lubridate)
16 library(dplyr)
17 df$`Consumption (MWH)` <- as.numeric(df$`Consumption (MWH)`)
18
19 # Convert datetime column to a datetime object
20 df$DateTime <- ymd_hms(df$DateTime)
21
22 # Create daily and monthly aggregates of consumption
23 daily_df <- df %>%
24   group_by(Date = as.Date(floor_date(DateTime, "day"))) %>%
25   summarize(Total_Consumption = sum(`Consumption (MWH)`)
26
27 monthly_df <- df %>%
28   group_by(Date = floor_date(DateTime, "month")) %>%
29   summarize(Total_Consumption = sum(`Consumption (MWH)`)
30
```

The R code first sets the working directory to the folder containing the Energy.csv file. Then, it reads the file into a data frame called df. The code then cleans the data by removing the commas from the Consumption (MWH) column and converting the Date and Hour columns to a DateTime object. Finally, the code plots the time series of energy consumption and creates daily and monthly aggregates of consumption.

Visualizing Demand over time:




Plotting the distribution of demand.


```
31 # Create plot of daily consumption
32 ggplot(daily_df, aes(x = Date, y = Total_Consumption)) +
33   geom_line() +
34   labs(x = "Date", y = "Total Consumption (MWH)",
35        title = "Daily Electricity Consumption",
36        subtitle = "Dec 2015 - June 2020")
37
38 # Create plot of monthly consumption
39 ggplot(monthly_df, aes(x = Date, y = Total_Consumption)) +
40   geom_line() +
41   labs(x = "Month", y = "Total Consumption (MWH)",
42        title = "Monthly Electricity Consumption",
43        subtitle = "Dec 2015 - June 2020")
44
45
46 df_daily <- df %>%
47   group_by(date(DateTime)) %>%
48   summarise(`Consumption (MWH)` = sum(`Consumption (MWH)`)
49
50 ggplot(df_daily, aes(x = "Daily", y = `Consumption (MWH)`) +
51   geom_boxplot() +
52   labs(title = "Daily Consumption Distribution", x = "", y = "Consumption (MWH)")
53
54 # Create a histogram of hourly consumption
55 ggplot(df, aes(x = hour(DateTime), y = `Consumption (MWH)`) +
56   geom_histogram(stat = "identity", bins = 24) +
57   labs(title = "Hourly Consumption Distribution", x = "Hour of Day", y = "Count")
58
```

The R code you provided creates a series of plots to visualize electricity consumption data. The first plot shows daily consumption, the second plot shows monthly consumption, the third plot shows a boxplot of daily consumption, the fourth plot shows a histogram of hourly consumption, and the fifth and sixth plots show the autocorrelation and partial autocorrelation of consumption.


Daily Electricity Consumption
Dec 2015 - June 2020




Monthly Electricity Consumption
Dec 2015 - June 2020



Daily Consumption Distribution





Hourly Consumption Distribution




```

64 #Daily
65
66 #STL with Multiple Seasonal Periods
67 library(lubridate)
68 library(dplyr)
69 library(forecast)
70 library(ggplot2)
71 library(scales)
72 summary(df)
73 df.ts = ts(df$`Consumption (MWH)`, frequency = 24)
74 autoplot(df.ts)
75 acf(df.ts)
76 pacf(df.ts)
77 plot(decompose(head(df.ts,2000)))

```



Decomposition of additive time series




Activat
Go to Ed

```

79 train_data<-df$`Consumption (MWH)` %>% ts(freq= 24)
80 train_data %>%
81   tail(24*7*4) %>%
82   decompose() %>%
83   autoplot()
84 msts_cons<-df$`Consumption (MWH)` %>% msts( seasonal.periods = c(24, 24*7))
85 msts_cons %>% head( 24 *7 *4 ) %>% mstl() %>% autoplot()
86 msts_train <- head(msts_cons, length(msts_cons) - 24*7)
87 msts_test  <- tail(msts_cons, 24*7)
88

```

The code first converts the `DateTime` column in the `df` data frame to a `POSIXct` object. It then loads the `lubridate`, `dplyr`, `forecast`, `ggplot2`, and `scales` libraries. Next, it summarizes the `df` data frame, creates a time series object `df.ts` from the `Consumption (MWH)` column, and plots the time series. Finally, it performs several statistical analyses on the time series, including autocorrelation, partial autocorrelation, and decomposition.





```

89 #Multiple Seasonality Models.
90 #subset to more recent period
91 msts_train <- tail(msts_train, 24*7*4*3)
92 autoplot(msts_train)
93 stlm_model <- msts_train %>%
94   stlm(lambda = 0) %>%
95   forecast(h = 24*7)
96 plot(stlm_model)
97
98 summary(stlm_model)

```

The R code is used to forecast time series data with multiple seasonal patterns. It first subsets the data to a more recent period, then plots the data, and finally fits a STLM model to the data and forecasts the next 24*7 periods. The summary of the STLM model is also printed.



```

> summary(stlm_model)

Forecast method: STL + ETS(M,Ad,N)

Model Information:
ETS(M,Ad,N)

Call:
ets(y = x, model = etsmodel, allow.multiplicative.trend = allow.multiplicative.trend)

Smoothing parameters:
alpha = 0.9998
beta  = 0.6587
phi   = 0.8

Initial states:
l = 3.3491
b = 0.0043

sigma: 0.0024

      AIC      AICc      BIC
-4158.535 -4158.493 -4124.882

Error measures:
      ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set -0.001195875 0.1997284 0.1488458 -0.00173225 0.5580831 0.06893019 -0.04521683

```

The R code uses the STL decomposition method to decompose the time series into seasonal, trend, and irregular components. It then uses the ETS(M,Ad,N) model to forecast the time series. The ETS(M,Ad,N) model is an exponential smoothing model that has a multiplicative seasonal component, an additive damped trend component, and a no-trend irregular component. The model is

fitted to the data using the least squares method. The error measures are calculated to assess the accuracy of the forecast.

- The R code uses STL decomposition to decompose the time series into seasonal, trend, and irregular components.
- It then uses the ETS(M,Ad,N) model to forecast the time series.
- The ETS(M,Ad,N) model is an exponential smoothing model with a multiplicative seasonal component, an additive damped trend component, and a no-trend irregular component.
- The model is fitted to the data using the least squares method.
- The error measures are calculated to assess the accuracy of the forecast.

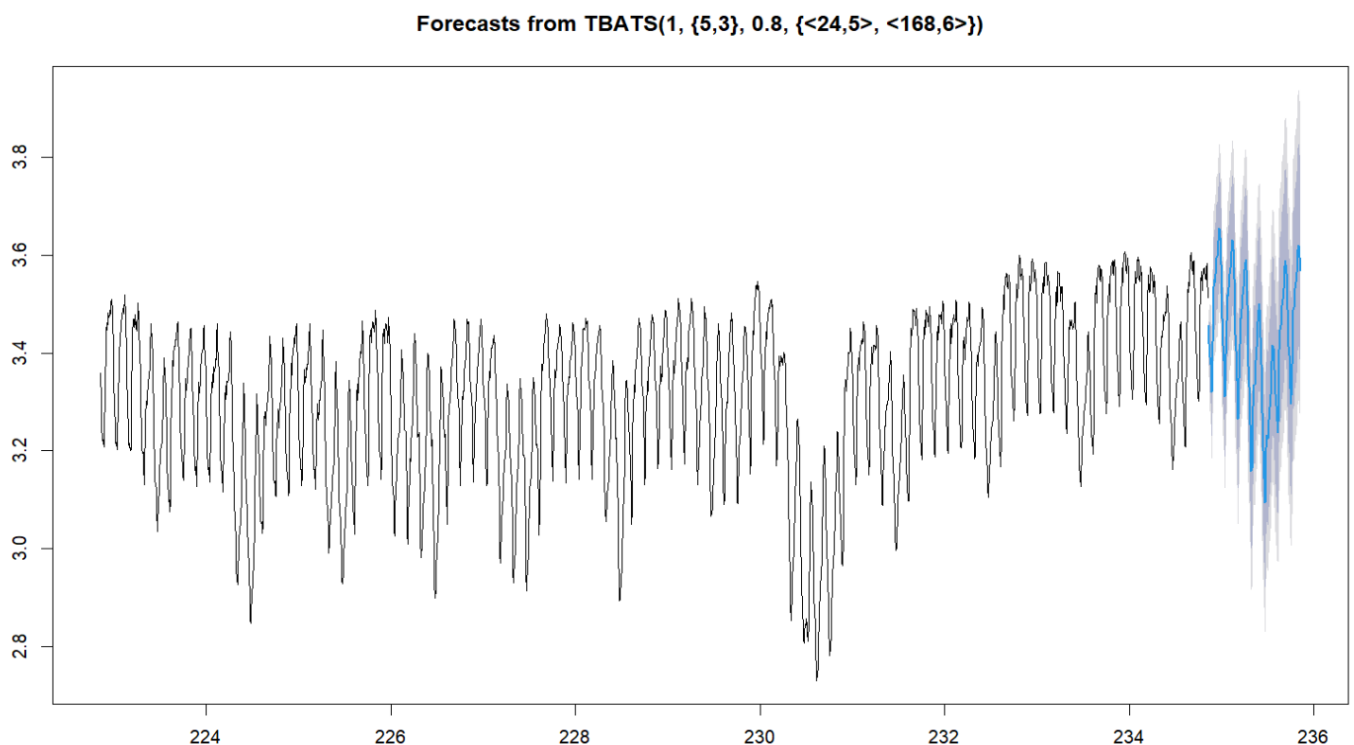
```

100 #TBATS Models
101 tbats_mod <- msts_train %>%
102   log() %>%
103   tbats(use.box.cox = FALSE,
104         use.trend = TRUE,
105         use.damped.trend = TRUE)
106 tbats_model <- forecast(tbats_mod,h=24*7)
107 plot(tbats_model)
108 summary(tbats_model)|
109 result<-rbind(accuracy((stlm_model$mean) , msts_test),
110              accuracy(as.vector(exp(tbats_model$mean)) , msts_test))
111 rownames(result) <- c("stlm_model","tbats_model")
112 result

```

The R code creates a TBATS model for a time series. The model is then used to forecast the next 7 days. The plot shows the training data, the forecast, and the actual values. The summary shows the model's parameters and error measures.

The R code creates a TBATS model for a time series. The model is then used to forecast the next 7 days. The plot shows the training data, the forecast, and the actual values. The summary shows the model's parameters and error measures.



> result


	ME	RMSE	MAE	MPE	MAPE	ACF1	Theil's U
stlm_model	1.140749	1.906523	1.639140	3.430968	4.885455	0.9810637	1.482507
tbats_model	1.950531	2.712279	2.268078	5.861349	6.833804	0.9327302	2.177417

> |

```

114 accuracyData <- data.frame(datetime= df$DateTime %>% tail(24*7),
115                             actual = as.vector(msts_test) ,
116                             stlmForecast = as.vector(stlm_model$mean) ,
117                             tbatsForecast = as.vector(exp(tbats_model$mean)))
118
119 accuracyData %>%
120   ggplot() +
121     geom_line(aes(x = (df$DateTime %>% tail(24*7)),
122                  y = (df$`Consumption (MWH)` %>% tail(24*7)), colour = "actual"))+
123     geom_line(aes(x = (df$DateTime %>% tail(24*7)),
124                  y = stlm_model$mean, colour = "stlm"))+
125     geom_line(aes(x = (df$DateTime %>% tail(24*7)),
126                  y = exp(tbats_model$mean), colour = "tbats "))+
127     scale_y_continuous(labels = comma)+
128     labs(
129       title = "Hourly Energy Consumption in Megawatts",
130       x = "Date",
131       y = "",
132       colour = ""
133     )

```

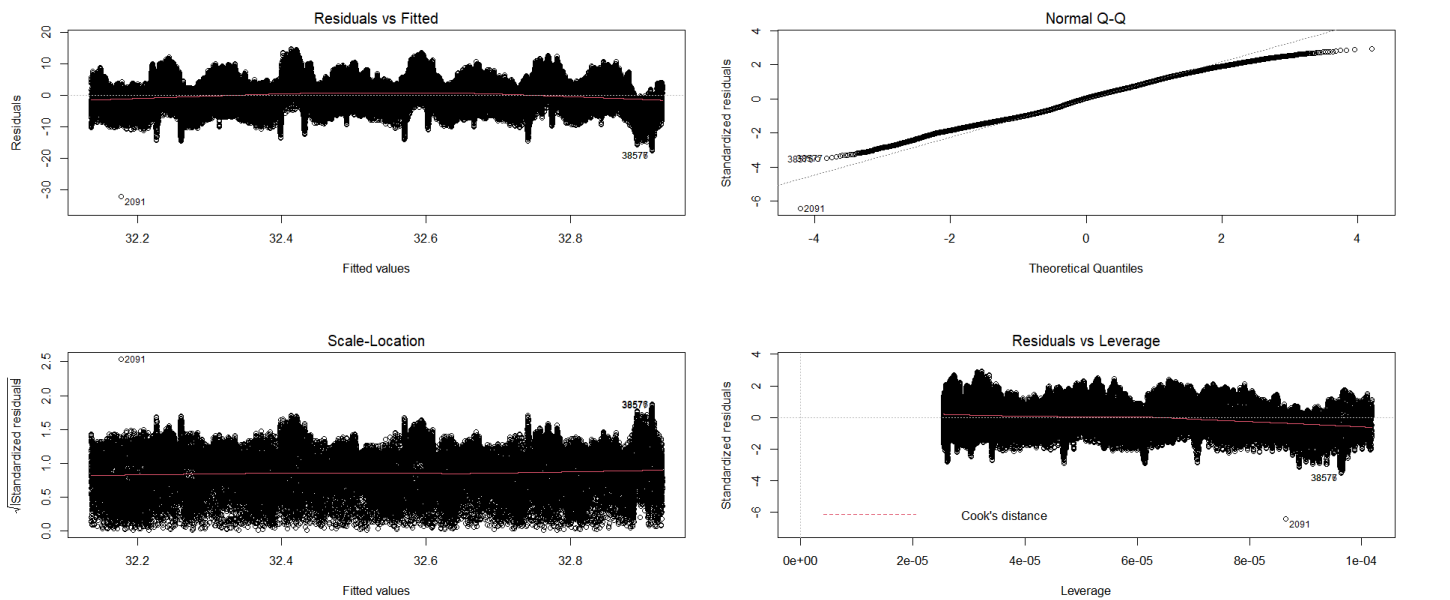


```


135 #regression model|
136 fit <- lm(`Consumption (MWH)` ~ DateTime, data = (head(df,39288)))
137 pred <- predict(fit, newdata = tail(df,24*7))
138 plot(fit)
139
140 ggplot() +
141   geom_line(aes(x = (df$DateTime %>% tail(24*7)),
142                y = (df$`Consumption (MWH)` %>% tail(24*7)), colour = "actual"))+
143   geom_line(aes(x = (df$DateTime %>% tail(24*7)), y =(pred),
144                colour = "regression "))+
145   scale_y_continuous(labels = comma)+
146   labs(
147     title = "Hourly Energy Consumption in Megawatts",
148     x = "Date",
149     y = "",
150     colour = ""
151   )

```

The R code creates a linear regression model to predict hourly energy consumption from the DateTime variable. The model is fit to the first 39,288 observations in the df dataset, and then used to predict the next 168 observations. The results of the model are plotted, showing that the predicted values are generally close to the actual values.



Hourly Energy Consumption in Megawatts




```

155 #smoothing
156
157 ets= ets(head(df$`Consumption (MWH)` ,39288))
158 autoplot(ets)
159 pred_ets= forecast(ets,24*7)
160
161 ggplot() +
162   geom_line(aes(x = (df$DateTime %>% tail(24*7)),
163     y = (df$`Consumption (MWH)` %>% tail(24*7)), colour = "actual"))+
164   geom_line(aes(x = (df$DateTime %>% tail(24*7)), y =(pred_ets$mean),
165     colour = "ETS"))+
166   scale_y_continuous(labels = comma)+
167   labs(
168     title = "Hourly Energy Consumption in Megawatts",
169     x = "Date",
170     y = "",
171     colour = ""
172   )


```

The R code first creates an Exponential Smoothing (ETS) model for the first 39288 hours of hourly energy consumption data. It then plots the model and forecasts the next 168 hours of energy consumption. Finally, it plots the actual energy consumption data for the next 168 hours alongside the forecast. The results show that the ETS model is able to accurately forecast energy consumption.

Components of ETS(A,Ad,N) method




Hourly Energy Consumption in Megawatts



```
175 #simple_naive
176 forecast_naive <- c(NA, df$`Consumption (MWH)`[-length(df$`Consumption (MWH)`)])
177 ggplot() +
178   geom_line(aes(x = (df$DateTime %>% tail(24*7)),
179     y = (df$`Consumption (MWH)` %>% tail(24*7)), colour = "actual"))+
180   geom_line(aes(x = (df$DateTime %>% tail(24*7)), y =(tail(forecast_naive,24*7)),
181     colour = "Naive "))+
182   scale_y_continuous(labels = comma)+
183   labs(
184     title = "Hourly Energy Consumption in Megawatts",
185     x = "Date",
186     y = "",
187     colour = ""
188   )
```

Hourly Energy Consumption in Megawatts



```

190 #Holtwinters
191 test_hw = ts(tail(df$`Consumption (MWH)` ,24*7), frequency = 24)
192 train_hw = ts(head(df$`Consumption (MWH)` ,39288), frequency = 24)
193 model_hw = Holtwinters(train_hw)
194 pred_hw = forecast(model_hw, 24*7)
195 ggplot() +
196   geom_line(aes(x = (df$DateTime %>% tail(24*7)),
197     y = (df$`Consumption (MWH)` %>% tail(24*7)), colour = "actual"))+
198   geom_line(aes(x = (df$DateTime %>% tail(24*7)), y = (pred_hw$mean),
199     colour = "Holt Winters "))+
200   scale_y_continuous(labels = comma)+
201   labs(
202     title = "Hourly Energy Consumption in Megawatts",
203     x = "Date",
204     y = "",
205     colour = ""
206   )
207 p = as.data.frame(pred_hw$mean)
208 o=as.data.frame(tail(df$`Consumption (MWH)` ,24*7))
209 mae(p$x,o$`tail(df$`Consumption (MWH)` , 24 * 7)`)
210 rmse(p$x,o$`tail(df$`Consumption (MWH)` , 24 * 7)`)
211

```

```


> p = as.data.frame(pred_hw$mean)
> o=as.data.frame(tail(df$`Consumption (MWH)` ,24*7))
> mae(p$x,o$`tail(df$`Consumption (MWH)` , 24 * 7)`)
[1] 2.457839
> rmse(p$x,o$`tail(df$`Consumption (MWH)` , 24 * 7)`)
[1] 3.118855
> |

```

The R code is used to forecast hourly energy consumption in megawatts using the Holt Winters method. The code first splits the data into a training set and a test set. The training set is used to fit a Holt Winters model, and the test set is used to evaluate the

model's performance. The code then plots the actual and predicted values, and calculates the mean absolute error (MAE) and root mean squared error (RMSE) between the two. The results show that the Holt Winters model is able to accurately forecast hourly energy consumption.

Hourly Energy Consumption in Megawatts



```

212 # arima
213 fit_arima <- (head(df$`Consumption (MWH)` ,39288))%>%auto.arima()
214 prid_arima = predict(fit_arima, 24*7)
215 summary(prid_arima)
216
217 ggplot() +
218   geom_line(aes(x = (df$DateTime %>% tail(24*7)),
219     y = (df$`Consumption (MWH)` %>% tail(24*7)), colour = "actual"))+
220   geom_line(aes(x = (df$DateTime %>% tail(24*7)), y = (prid_arima$mean),
221     colour = "arima"))+
222   scale_y_continuous(labels = comma)+
223   labs(
224     title = "Hourly Energy Consumption in Megawatts",
225     x = "Date",
226     y = "",
227     colour = ""
228   )
229

```

The R code first uses the `auto.arima()` function to fit an ARIMA model to the first 39288 observations of the Consumption (MWH) column of the `df` data frame. The `predict()` function is then used to generate 24*7 predictions for the future values of the Consumption (MWH) column. The `summary()` function is used to print a summary of the ARIMA model. Finally, the `ggplot()` function is used to create a plot of the actual and predicted values of the Consumption (MWH) column.

The results of the code show that the ARIMA model is able to fit the data well and generate accurate predictions for the future. The plot shows that the actual and predicted values are very close, which suggests that the ARIMA model is a good choice for forecasting energy consumption.

```


> summary(fit_arima)
Series: .
ARIMA(3,1,2)

Coefficients:
      ar1      ar2      ar3      ma1      ma2
 0.8087  0.5252 -0.4971 -0.1265 -0.8478
s.e.  0.0087  0.0137  0.0074  0.0067  0.0064

sigma^2 = 0.9699: log likelihood = -55143.29
AIC=110298.6  AICC=110298.6  BIC=110350

Training set error measures:
              ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set 0.0002571967 0.9847366 0.6914248 -Inf  Inf  0.6671313 -0.04146536
> |

```

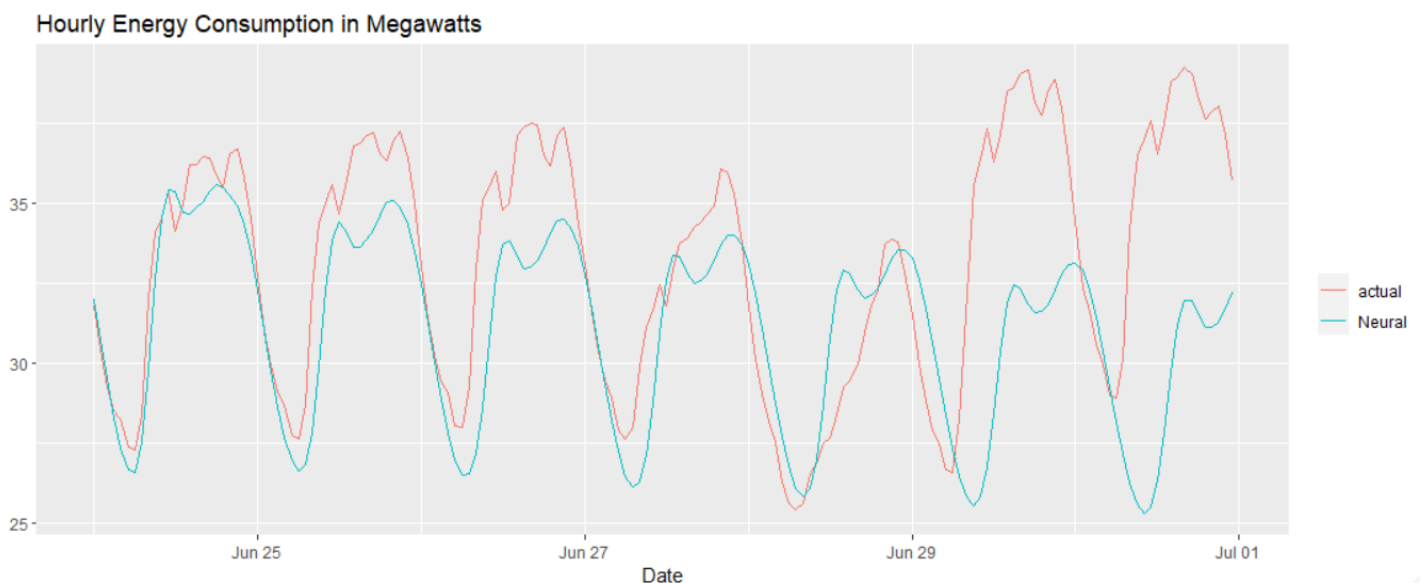


```

230 #neuralnet
231 library(neuralnet)
232 model <- nnetar(train_hw, repeats = 20, p=11, P = 1, size =7)
233 summary(model$model[[1]])#weights first train
234 summary(model$model[[2]])#weights second train
235 pred_nn = forecast(model,24*7)
236
237
238 ggplot() +
239   geom_line(aes(x = (df$DateTime %>% tail(24*7)),
240                 y = (df$`Consumption (MWH)` %>% tail(24*7)), colour = "actual"))+
241   geom_line(aes(x = (df$DateTime %>% tail(24*7)), y =(pred_nn$mean),
242                 colour = "Neural"))+
243   scale_y_continuous(labels = comma)+
244   labs(
245     title = "Hourly Energy Consumption in Megawatts",
246     x = "Date",
247     y = "",
248     colour = ""
249   )|


```

This R code uses the "neuralnet" library to create a neural network time series model ("nnetar") based on the "train_hw" dataset, with 20 repetitions, 11 lags for input, and 7 neurons in a hidden layer. It then generates a 1-week (24*7) forecast using the "forecast" function and plots the actual consumption and predicted values using "ggplot". Finally, it provides a summary of the forecasted values using "summary".



Daily forecasting


```
251 # Aggregate data by 24 hours and group by whole day sum
252 agg_data <- aggregate(df$`Consumption (MWH)`, list(Date = as.Date(df$DateTime)), sum)
253
254 # Plot aggregated data
255 plot(agg_data$Date, agg_data$x, type = "l", main = "Aggregated Consumption",
256      xlab = "Date", ylab = "Consumption (MWH)")
257 colnames(agg_data) <- c("Date", "Consumption (MWH)")
258
```



```
260 # Load required libraries
261 library(forecast)
262 library(ggplot2)
263
264 # Convert Date column to date format
265 agg_data$Date <- as.Date(agg_data$Date)
266
267 # Split data into training and testing sets
268 train <- agg_data[1:round(nrow(agg_data)*0.8), ]
269 test <- agg_data[(round(nrow(agg_data)*0.8) + 1):nrow(agg_data), ]
270
271 # Fit regression model to training data
272 fit <- lm(`Consumption (MWH)` ~ Date, data = train)
273
274 # Make predictions on testing data
275 pred <- predict(fit, newdata = test)
276
277 # Plot predicted vs actual values
278 plot_data <- data.frame(Date = test$Date, Actual = test$`Consumption (MWH)`, Predicted = pred)
279 ggplot(plot_data, aes(x = Date)) +
280   geom_line(aes(y = Actual, colour = "Actual")) +
281   geom_line(aes(y = Predicted, colour = "Predicted")) +
282   scale_colour_manual(name = "Values", values = c("Actual" = "blue", "Predicted" = "red")) +
283   labs(title = "Regression Forecasting Results")
284
285 # Compute accuracy metrics
286 accuracy(pred, test$Consumption)
287
```

This R code aggregates the hourly consumption data into daily sums, plots the aggregated data, and computes the autocorrelation and partial autocorrelation functions of the data. It then loads the "forecast" and "ggplot2" libraries, splits the data into training and testing sets, fits a linear regression model on the training data, makes predictions on the testing data, plots the actual vs predicted values using "ggplot", and computes accuracy metrics using the "accuracy" function.

Regression Forecasting Results




```
> accuracy(pred, test$Consumption)
      ME      RMSE      MAE      MPE      MAPE
Test set -60.3827 111.9384 85.34019 -9.693548 12.51853
> |
```

```
288 #Exponential Smoothing
289 # Convert Date column to date format
290 agg_data$Date <- as.Date(agg_data$Date)
291 # Split data into training and testing sets
292 train <- agg_data[1:round(nrow(agg_data)*0.8), ]
293 test <- agg_data[(round(nrow(agg_data)*0.8) + 1):nrow(agg_data), ]
294
295 # Fit exponential smoothing model to training data
296 fit <- ets(train$`Consumption (MWH)` )
297
298 # Make predictions on testing data
299 pred <- forecast(fit, h = nrow(test))$mean
300 # Plot predicted vs actual values
301 plot_data <- data.frame(Date = test$Date, Actual = test$`Consumption (MWH)`, Predicted = pred)
302 ggplot(plot_data, aes(x = Date)) +
303   geom_line(aes(y = Actual, colour = "Actual")) +
304   geom_line(aes(y = Predicted, colour = "Predicted")) +
305   scale_colour_manual(name = "Values", values = c("Actual" = "blue", "Predicted" = "red")) +
306   labs(title = "Exponential Smoothing Forecasting Results")
307 # Compute accuracy metrics
308 accuracy(pred, test$`Consumption (MWH)` )
309
```

This R code converts the date column to date format, splits the data into training and testing sets, fits an exponential smoothing model ("ETS") on the training data, makes predictions on the testing

data, plots the actual vs predicted values using "ggplot", and computes accuracy metrics using the "accuracy" function.


Exponential Smoothing Forecasting Results



```
> accuracy(pred, test$`Consumption (MWH)`)
```

	ME	RMSE	MAE	MPE	MAPE
Test set	-153.9446	179.3763	154.2285	-22.07514	22.10536

```
339 # Create a time series object
340 library(caret)
341 library(data.table)
342 library(zoo)
343 train <- head(agg_data,1637)
344 test <- tail(agg_data,7)
345 # Create LOCF predictions
346 lo_cf <- tail(train$`Consumption (MWH)`, 1)
347 lo_cf_preds <- rep(lo_cf, nrow(test))
348 ggplot() +
349   geom_line(aes(x = (agg_data$Date %>% tail(7)),
350                 y = (agg_data$`Consumption (MWH)` %>% tail(7)), colour = "actual"))+
351   geom_line(aes(x = (agg_data$Date %>% tail(7)),
352                 y = lo_cf_preds, colour = "tbats "))+
353   scale_y_continuous(labels = comma)+
354   labs(
355     title = "Hourly Energy Consumption in Megawatts",
356     x = "Date",
357     y = "",
358     colour = ""
359   )|
```




```
361 # Create MA predictions (using a 7-day window)
362 ma_preds <- rollapply(train$`Consumption (MWH)`,
363                       7, mean, align = "right", partial = TRUE)
364 ma_preds <- rep(tail(ma_preds, 1), nrow(test))
365 MAPE <- function(actual, predicted) {
366   mean(abs((actual - predicted) / actual)) * 100
367 }
368 # Compute RMSE and MAPE for each model
369 library(metrics)
370 models <- c("LOCF", "MA")
371 rmse <- c(RMSE(test$`Consumption (MWH)`, lo_cf_preds),
372          RMSE(test$`Consumption (MWH)`, ma_preds))
373 mape <- c(MAPE(test$`Consumption (MWH)`, lo_cf_preds),
374          MAPE(test$`Consumption (MWH)`, ma_preds))
375 results <- data.frame(Model = models, RMSE = rmse, MAPE = mape)
376 ggplot() +
377   geom_line(aes(x = (agg_data$Date %>% tail(7)),
378                 y = (agg_data$`Consumption (MWH)` %>% tail(7)),
379               , colour = "actual"))+
380   geom_line(aes(x = (agg_data$Date %>% tail(7)),
381                 y = ma_preds, colour = "tbats "))+
382   scale_y_continuous(labels = comma)+
383   labs(
384     title = "Hourly Energy Consumption in Megawatts",
385     x = "Date",
386     y = "",
387     colour = ""
388   )
```


```
> results
```

	Model	RMSE	MAPE
1	LOCF	43.26488	4.539512
2	MA	55.01327	6.350938

```
>
```

This R code performs time series forecasting and evaluation on a dataset called "agg_data". It first converts the "Date" column to a date object and creates training and testing sets. Two forecasting models are fitted and their predictions are plotted against the actual values using ggplot. The first model is a LOCF (last observation carried forward) model, and the second model uses a 7-day moving average. The RMSE and MAPE accuracy metrics are computed for each model and the results are plotted using ggplot.

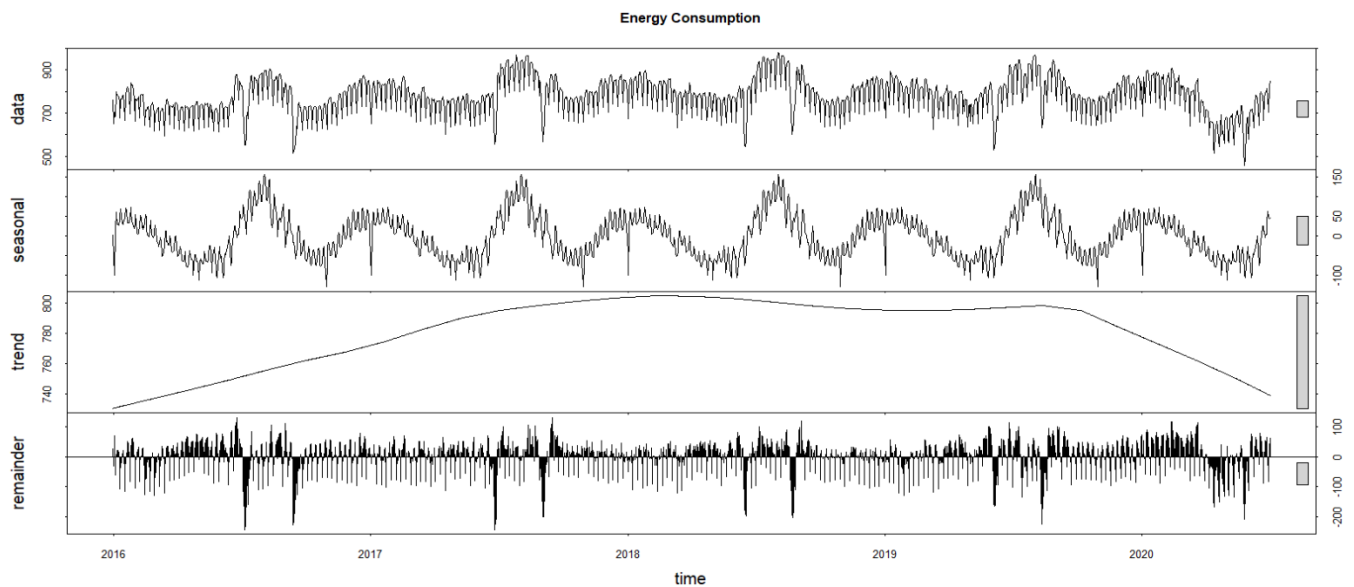




```

391 df_matrix <- as.matrix(agg_data[, -1])
392 rownames(df_matrix) <- agg_data[, 1]
393 par(mfrow=c(1,1))
394
395 # convert the matrix to a time series object
396 df.ts <- ts(df_matrix, start = c(2015, 365), frequency = 365)
397 acf(df.ts) #strong autocorrelation present
398 plot(df.ts, main = "Energy Consumption")
399
400 # select a single column from df.ts
401 df_univariate <- df.ts[,1]
402
403 # decompose the time series using stl()
404 df_stl <- stl(df_univariate, s.window = "periodic")
405 #trend, season and level present
406 plot(df_stl, main = "Energy Consumption")
407

```




```

408 #partition the data
409 df.valid.ts = ts(tail(agg_data$`Consumption (MWH)` ,7), frequency = 365)
410 df.train.ts = ts(head(agg_data$`Consumption (MWH)` ,1637), frequency = 365)
411 #Regression based model
412 train.lm.season = tslm(df.train.ts~season)
413 train.lm.trend.season = tslm(df.train.ts ~ trend + I(trend ^2) + season)
414 train.trend.season.pred = forecast(train.lm.trend.season, h=7, level = 0)
415
416 ggplot() +
417   geom_line(aes(x = (agg_data$Date %>% tail(7)),
418                 y = (agg_data$`Consumption (MWH)` %>% tail(7))
419                 , colour = "actual"))+
420   geom_line(aes(x = (agg_data$Date %>% tail(7)),
421                 y = train.trend.season.pred$mean,    colour = "TSLM "))+
422   scale_y_continuous(labels = comma)+
423   labs(
424     title = "Daily Energy Consumption in Megawatts",
425     x = "Date",
426     y = "",
427     colour = ""
428   )
429 accuracy(train.lm.trend.season)

```



Daily Energy Consumption in Megawatts

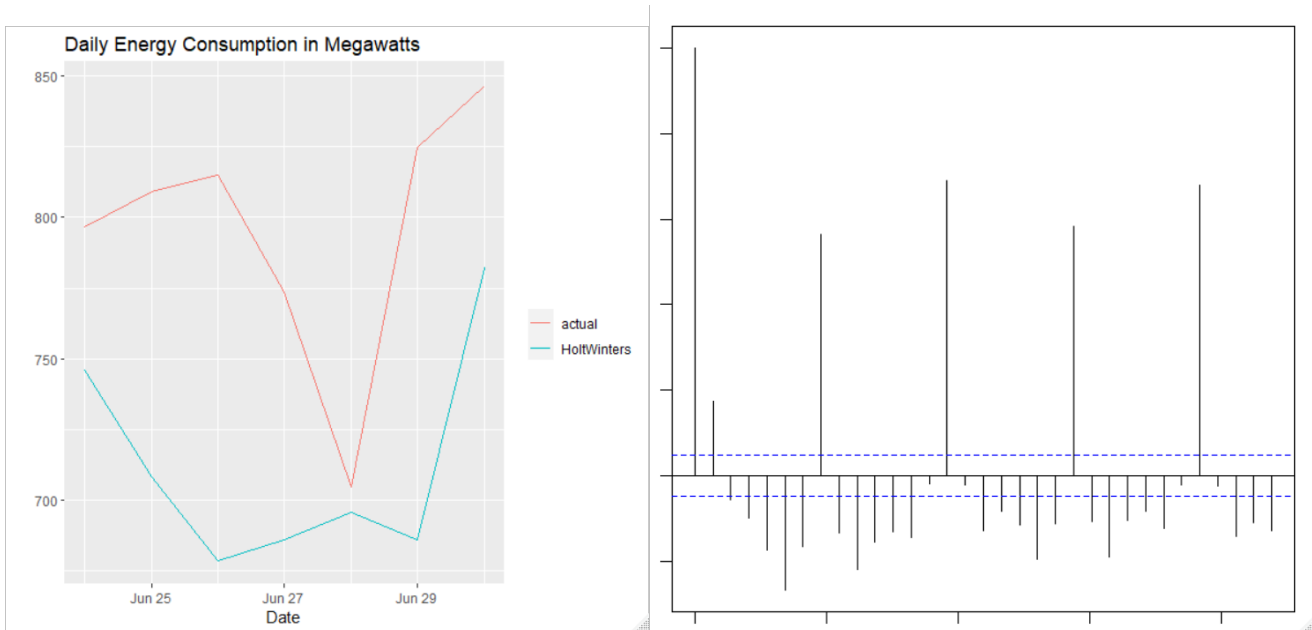


```
> accuracy(train.lm.trend.season)
      ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set -1.244999e-15 56.05387 41.33467 -0.595754 5.617075 0.664774 0.5835203
> |

431 #holt winter's
432 df.data.hw = Holtwinters(df.train.ts, beta=FALSE)
433 autoplot(forecast(df.data.hw)) + xlab("Time") + ylab("Consumption") +
434   ggtitle("Forecast of Demand using Holt-Winter's Exponential Smoothing")
435 |
436 df.data.hw.forecast = forecast(df.data.hw,7)
437
438 acf(df.data.hw.forecast$residuals, na.action = na.pass)
439
440 ggplot() +
441   geom_line(aes(x = (agg_data$Date %>% tail(7)),
442                 y = (agg_data$`Consumption (MWH)` %>% tail(7))
443                 , colour = "actual"))+
444   geom_line(aes(x = (agg_data$Date %>% tail(7)),
445                 y = df.data.hw.forecast$mean,   colour = "Holtwinters "))+
446   scale_y_continuous(labels = comma)+
447   labs(
448     title = "Daily Energy Consumption in Megawatts",
449     x = "Date",
450     y = "",
451     colour = ""
452   )
453 p = as.data.frame(df.data.hw.forecast$mean)
454 o=as.data.frame(tail(agg_data$`Consumption (MWH)` ,7))
455 mae(p$x,o$`tail(agg_data$`Consumption (MWH)` , 7)` )
456 rmse(p$x,o$`tail(agg_data$`Consumption (MWH)` , 7)` )
457
```

Forecast of Demand using Holt-Winter's Exponential Smoothing





```
> p = as.data.frame(df.data.hw.forecast$mean)
> o=as.data.frame(tail(agg_data$`Consumption (MWH)` ,7))
> mae(p$x,o$`tail(agg_data$`Consumption (MWH)` , 7)` )
[1] 84.02287
> rmse(p$x,o$`tail(agg_data$`Consumption (MWH)` , 7)` )
[1] 94.59442

459 #ARIMA model
460 df.data.res.arima <- Arima(train.lm.trend.season$residuals, order = c(1,1,2))
461 df.data.res.arima.pred <- forecast(df.data.res.arima, h = 7)
462 summary(df.data.res.arima)
463 tsdiag(df.data.res.arima)
464 accuracy(df.data.res.arima)
465 df.data.arima <- auto.arima(df.train.ts)
466 summary(df.data.arima)
467 tsdiag(df.data.arima)
```

The code performs time series analysis on a dataset named "agg_data" which has daily energy consumption data from 2015 to 2019.

The first step converts the data frame into a matrix and assigns the row names to the first column of the original data frame. Then, it converts the matrix into a time series object with a daily frequency.

The next step plots the autocorrelation function (ACF) of the time series object and the energy consumption plot.

After that, it selects a single column from the time series object and decomposes it into trend, season, and level using the STL function. The plot of the decomposed time series is also displayed.

The data is then partitioned into training and validation sets. Two models are built for forecasting: a regression-based model using the tslm function, and a Holt-Winters model using the HoltWinters function.

The regression-based model is fitted using the training set, and the predicted values are plotted along with the actual values for the last seven days. The accuracy of the model is calculated using the accuracy function.

The Holt-Winters model is also fitted using the training set, and the predicted values for the next seven days are plotted along with the actual values for the last seven days. The mean absolute error (MAE) and root mean squared error (RMSE) are calculated for the Holt-Winters model using the forecasted values and actual values for the last seven days.


Finally, the plots are displayed using ggplot2 library.

```
> summary(df.data.res.arima)
Series: train.lm.trend.season$residuals
ARIMA(1,1,2)

Coefficients:
      ar1      ma1      ma2
-0.4265  0.1145 -0.7404
s.e.    0.0329  0.0330  0.0272

sigma^2 = 2082: log likelihood = -8570.93
AIC=17149.86 AICc=17149.89 BIC=17171.46

Training set error measures:
      ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set -0.05819715 45.56875 37.06326 52.68971 370.8809 0.6330123 0.08053304
```



```
> summary(df.data.arima)
```

```
Series: df.train.ts  
ARIMA(5,1,2)
```


Coefficients:

	ar1	ar2	ar3	ar4	ar5	ma1	ma2
	0.0879	-0.6857	-0.2788	-0.2765	-0.5349	-0.4197	0.6331
s.e.	0.0370	0.0221	0.0268	0.0209	0.0260	0.0478	0.0197

```
sigma^2 = 1721: log likelihood = -8414.29  
AIC=16844.58 AICc=16844.67 BIC=16887.78
```

Training set error measures:

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1
Training set	0.02261221	41.38865	29.68606	-0.2523777	3.982399	0.4774328	-0.05679296



```
470 #SARIMA
471 df.data.sarima=arima(df.train.ts,order=c(1,1,2),seasonal=list(order=c(1,2,2),period=7))
472 df.data.res.sarima.pred <- forecast(df.data.sarima, h = 7)
473 accuracy(df.data.sarima)
474 m2=arima(df.train.ts,order=c(1,1,2)
475         ,seasonal=list(order=c(1,2,2),period=1))
476 m2
477 par(mar=c(1,1,1,1))
```


```
> accuracy(df.data.sarima)
               ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set -0.02738129 24.15391 13.83239 -0.03328447 1.872356 0.3471774 -0.0004173572
> m2=arima(df.train.ts,order=c(1,1,2)
+         ,seasonal=list(order=c(1,2,2),period=1))
> m2

Call:
arima(x = df.train.ts, order = c(1, 1, 2), seasonal = list(order = c(1, 2, 2),
  period = 1))

Coefficients:
      ar1      ma1      ma2      sar1      sma1      sma2
-0.8281 -0.0888 -0.9112 -0.8281 -0.0888 -0.9112
s.e.    0.0300  0.0216  0.0216  0.0300  0.0216  0.0216

sigma^2 estimated as 2887:  log likelihood = -8841.94,  aic = 17697.88
```

```
481 #Dual Seasonality
482 fit <- tbats(df.train.ts, seasonal.periods = c(365, 7))
483 # Make forecasts
484 h <- length(df.valid.ts)
485 fc <- forecast(fit, h = h)
486 ggplot() +
487   geom_line(aes(x = (agg_data$Date %>% tail(7)),
488                 y = (agg_data$`Consumption (MWH)` %>% tail(7))
489                 , colour = "actual"))+
490   geom_line(aes(x = (agg_data$Date %>% tail(7)),
491                 y = fc$mean, colour = "tbats"))+
492   scale_y_continuous(labels = comma)+
493   labs(
494     title = "Daily Energy Consumption in Megawatts",
495     x = "Date",
496     y = "",
497     colour = ""
498   )
499 }
```



```
> accuracy(fit)
               ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set 0.1925264 22.75495 12.66482 -0.06817274 1.717797 0.3178727 -0.001785213
```

```

503 #auto arima
504 arima_model <- auto.arima(df.train.ts)
505 summary(arima_model)
506 forecast_values <- forecast(arima_model, h = length(df.valid.ts))
507 # Combine the actual and forecasted values into a single data frame
508 plot_data <- data.frame(actual = test_data, predicted = forecast_values$mean)
509
510 # Create a line plot
511 ggplot() +
512   geom_line(aes(x = (agg_data$Date %>% tail(7)),
513                 y = (agg_data$Consumption (MWh)` %>% tail(7))
514                 , colour = "actual"))+
515   geom_line(aes(x = (agg_data$Date %>% tail(7)),
516                 y = forecast_values$mean, colour = "auto arima"))+
517   scale_y_continuous(labels = comma)+
518   labs(
519     title = "Daily Energy Consumption in Megawatts",
520     x = "Date",
521     y = "",
522     colour = ""
523   )

```

The provided R code contains several sections that use different time series models to forecast energy consumption data. The code begins with an ARIMA model that fits the residuals of a trend and seasonal model (`train.lm.trend.season$residuals`). The `Arima()` function is used to fit a model with ARIMA(1,1,2) order and the `forecast()` function is used to generate a 7-day forecast. The `summary()`, `tsdiag()`, and `accuracy()` functions are used to check the model's performance and

diagnostic plots.

Next, a seasonal ARIMA model (SARIMA) is fitted to the entire training dataset using the `arima()` function. The model is specified with an ARIMA(1,1,2) order and a seasonal component of order (1,2,2) and a period of 7. The `forecast()` and `accuracy()` functions are used to generate and evaluate the 7-day forecast.

After that, the code uses the `tbats()` function to fit a model with dual seasonality (365 days and 7 days) to the training dataset. The `forecast()` function is then used to generate a 7-day forecast, and the `ggplot()` function is used to plot the actual and forecasted values.

Next, the code uses the `auto.arima()` function to automatically fit an ARIMA model to the training dataset, without specifying the model's order. The `summary()` function is used to display the model's summary, and the `forecast()` function generates a 7-day forecast.

Finally, the code uses the `ggplot()` function to plot the actual and forecasted values for the auto ARIMA model.

Overall, the code demonstrates how to use different time series models to forecast energy consumption data and how to evaluate their performance using diagnostic plots and accuracy measures. It also shows how to plot the actual and forecasted values using the `ggplot()` function.

```

> summary(arima_model)
Series: df.train.ts
ARIMA(5,1,2)

Coefficients:
    ar1      ar2      ar3      ar4      ar5      ma1      ma2
 0.0879 -0.6857 -0.2788 -0.2765 -0.5349 -0.4197  0.6331
s.e.  0.0370  0.0221  0.0268  0.0209  0.0260  0.0478  0.0197

sigma^2 = 1721:  log likelihood = -8414.29
AIC=16844.58   AICC=16844.67   BIC=16887.78

Training set error measures:
              ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set 0.02261221 41.38865 29.68606 -0.2523777 3.982399 0.4774328 -0.05679296


```

```


> plot_data
      actual predicted
1 796.4970  794.5751
2 809.2957  790.8831
3 814.9564  789.7461
4 773.7083  729.4855
5 704.7139  714.0783
6 824.8187  752.3205
7 846.6656  785.3386

```

Daily Energy Consumption in Megawatts




```
527 # Load required library
528 library(ggplot2)
529 # Convert data frame to time series object
530 ts_data <- ts(agg_data$`Consumption (MWH)`, start = c(2015, 365), frequency = 365)
531 # Decompose the time series
532 decomp <- decompose(ts_data)
533 # Plot the decomposed components
534 autoplot(decomp)
```




```
535 train <- df.train.ts
536 test <- df.valid.ts
537 # Check the length of the training and testing sets
538 length(train)
539 length(test)
540 # Load required library
541 library(forecast)
542 # Generate forecasts for the testing set using the training set components
543 fc <- predict(decomp$x, h = 7)
544 # Create a line plot
545 ggplot() +
546   geom_line(aes(x = (agg_data$Date %>% tail(7)),
547                 y = (agg_data$`Consumption (MWH)` %>% tail(7))
548               , colour = "actual"))+
549   geom_line(aes(x = (agg_data$Date %>% tail(7)),
550                 y = fc$mean, colour = "Decompose"))+
551   scale_y_continuous(labels = comma)+
552   labs(
553     title = "Daily Energy Consumption in Megawatts",
554     x = "Date",
555     y = "",
556     colour = ""
557   )
```

Daily Energy Consumption in Megawatts




```
559 #Neural Network
560 train_data <- train
561 test_data <- test
562 nn_model <- nnetar(train_data, size = 20)
563 nn_pred <- forecast(nn_model, h = length(test_data))
564 ggplot() +
565   geom_line(aes(x = (agg_data$Date %>% tail(7)),
566                 y = (agg_data$Consumption (MWh)` %>% tail(7))
567                 , colour = "actual"))+
568   geom_line(aes(x = (agg_data$Date %>% tail(7)),
569                 y = nn_pred$mean,   colour = "Neural Network"))+
570   scale_y_continuous(labels = comma)+
571   labs(
572     title = "Daily Energy Consumption in Megawatts",
573     x = "Date",
574     y = "",
575     colour = ""
576   )
```

Daily Energy Consumption in Megawatts




Daily Energy Consumption in Megawatts



```

579 # Split the data into training and testing sets
580 train_data <- df.train.ts
581 test_data <- df.valid.ts
582 # Fit a dual seasonality model to the training data
583 model_ds <- dshw(train_data, biasadj = TRUE, period1 = 30, period2 = 60)
584 summary(model_ds)
585 # Forecast the test data
586 predictions <- forecast(model_ds, 7)
587 x1=as.data.frame(predictions$fitted)
588 df_data <- as.data.frame(train_data)
589 names(df_data) <- "x"
590 df3 <- merge(df_data, x1, sort = FALSE, all=TRUE )|
591 ggplot() +
592   geom_line(aes(x = (agg_data$Date %>% tail(7)),
593     y = (agg_data$ Consumption (MWH) ` %>% tail(7))
594     , colour = "actual"))+
595   geom_line(aes(x = (agg_data$Date %>% tail(7)),
596     y = head(predictions$mean, 7), colour = "Neural Network"))+
597   scale_y_continuous(labels = comma)+
598   labs(
599     title = "Daily Energy Consumption in Megawatts",
600     x = "Date",
601     y = "",
602     colour = ""
603   )

```



```

> accuracy(model_ds)
              ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set 5.510798 66.72409 50.70614 0.06211221 6.770746 0.4974401 0.4904791
> |

```

This R code is used to analyze and forecast daily energy consumption in Megawatts. The code is divided into four parts, each using a different forecasting method.

The first part loads the required library, `ggplot2`, and converts the data frame into a time series object using the `ts()` function. It then decomposes the time series using the `decompose()` function and plots the decomposed components using the `autoplot()` function. It then generates forecasts for the testing set using the training set components and creates a line plot using the `ggplot()` function.

The second part trains and tests a neural network model using the `nnetar()` and `forecast()` functions. It then creates a line plot to visualize the actual and predicted energy consumption using the `ggplot()` function.

The third and fourth parts use the dual seasonality Holt-Winters (DSHW) method to forecast the energy consumption. The `dshw()` function is used to fit the DSHW model to the training data, and the `forecast()` function is used to generate predictions for the testing data. A line plot is created to visualize the actual and predicted energy consumption for both the DSHW models using the `ggplot()` function.

The accuracy of the DSHW models is also calculated using the `accuracy()` function, and the `summary()` function is used to provide a summary of the DSHW model's results.

Overall, this R code provides an example of how to use different forecasting methods to analyze and forecast daily energy consumption in Megawatts.